

Understanding stacks and registers in the Sparc architecture(s)

The [Sparc](#) architecture from [Sun Microsystems](#) has some "interesting" characteristics. After having to deal with both compiler, interpreter, OS emulator, and OS porting issues for the Sparc, I decided to gather notes and documentation in one place. If there are any issues you don't find addressed by this page, or if you know of any similar Net resources, let me know. This document is limited to the V8 version of the architecture.

General Structure

Sparc has 32 general purpose integer registers visible to the program at any given time. Of these, 8 registers are *global* registers and 24 registers are in a register *window*. A window consists of three groups of 8 registers, the *out*, *local*, and *in* registers. See table 1. A Sparc [implementation](#) can have from 2 to 32 windows, thus varying the number of registers from 40 to 520. Most implementations have 7 or 8 windows. The variable number of registers is the principal reason for the Sparc being "scalable".

At any given time, only one window is visible, as determined by the *current window pointer* (CWP) which is part of the processor status register (PSR). This is a five bit value that can be decremented or incremented by the SAVE and RESTORE instructions, respectively. These instructions are generally executed on procedure call and return (respectively). The idea is that the *in* registers contain incoming parameters, the *local* register constitute scratch registers, the *out* registers contain outgoing parameters, and the *global* registers contain values that vary little between executions. The register windows overlap partially, thus the *out* registers become renamed by SAVE to become the *in* registers of the called procedure. Thus, the memory traffic is reduced when going up and down the procedure call. Since this is a frequent operation, performance is improved.

(That was the idea, anyway. The drawback is that upon interactions with the system the registers need to be flushed to the stack, necessitating a long sequence of writes to memory of data that is often mostly garbage. Register windows was a bad idea that was caused by simulation studies that considered only programs in isolation, as opposed to multitasking workloads, and by considering compilers with poor optimization. It also caused considerable problems in implementing high-end Sparc processors such as the SuperSparc, although more recent implementations have dealt effectively with the obstacles. Register windows is now part of the compatibility legacy and not easily removed from the architecture.)

Register Group	Mnemonic	Register Address
global	%g0-%g7	r[0]-r[7]
out	%o0-%o7	r[8]-r[15]
local	%l0-%l7	r[16]-r[23]
in	%i0-%i7	r[24]-r[31]

Table 1 - Visible Registers

The overlap of the registers is illustrated in figure 1. The figure shows an implementation with 8 windows, numbered 0 to 7 (labeled w0 to w7 in the figure).. Each window corresponds to 24 registers, 16 of which are shared with "neighboring" windows. The windows are arranged in a wrap-around manner, thus window number 0 borders window number 7. The common cause of

changing the current window, as pointed to by CWP, is the RESTORE and SAVE instructions, shown in the middle. Less common is the supervisor RETT instruction (return from trap) and the trap event (interrupt, exception, or TRAP instruction).

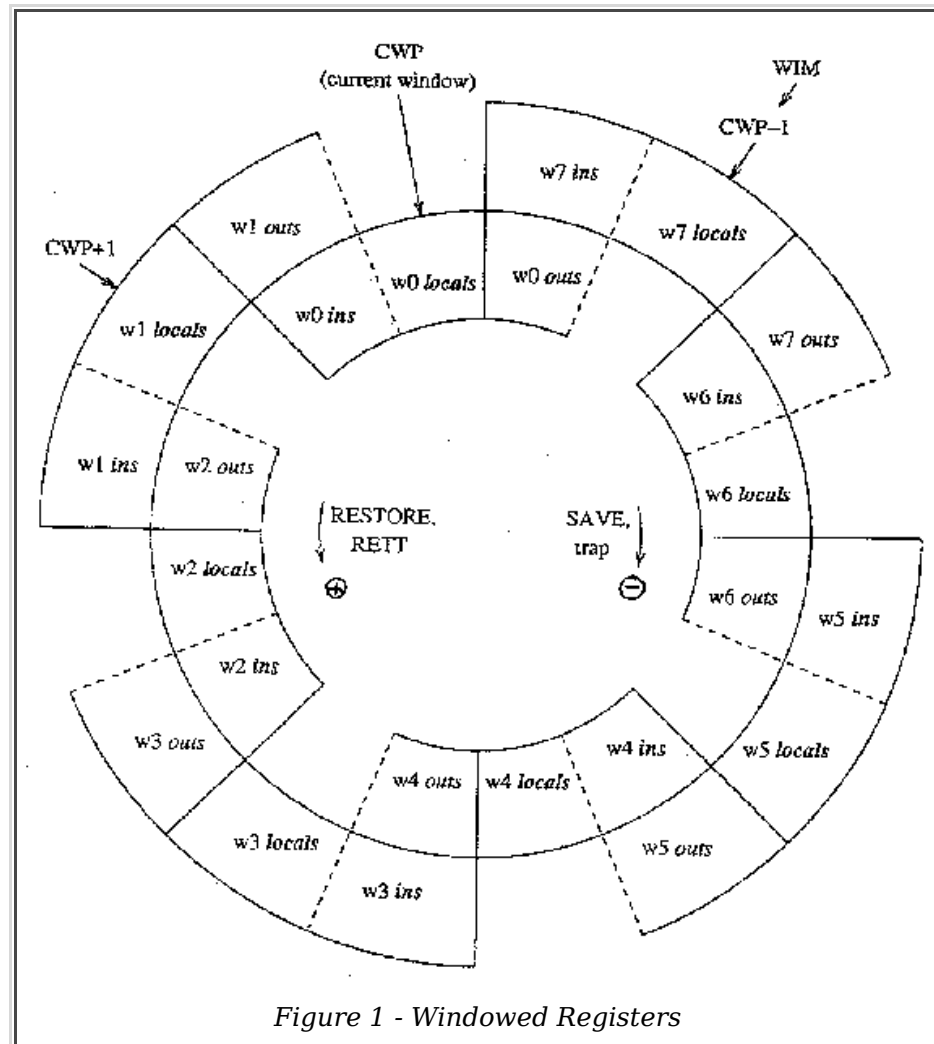


Figure 1 - Windowed Registers

The "WIM" register is also indicated in the top left of figure 1. The *window invalid mask* is a bit map of valid windows. It is generally used as a pointer, i.e. exactly one bit is set in the WIM register indicating which window is invalid (in the figure it's window 7). Register windows are generally used to support procedure calls, so they can be viewed as a cache of the stack contents. The WIM "pointer" indicates how many procedure calls in a row can be taken without writing out data to memory. In the figure, the capacity of the register windows is fully utilized. An additional call will thus exceed capacity, triggering a *window overflow* trap. At the other end, a *window underflow* trap occurs when the register window "cache" is empty and more data needs to be fetched from memory.

Register Semantics

The Sparc Architecture includes *recommended* software semantics. These are described in the architecture manual, the Sparc ABI (application binary interface) standard, and, unfortunately, in various other locations as well (including header files and compiler documentation).

Figure 2 shows a summary of register contents at any given time.

global	%g0	(r00)		always zero
	%g1	(r01)	[1]	temporary value
	%g2	(r02)	[2]	global 2
	%g3	(r03)	[2]	global 3
	%g4	(r04)	[2]	global 4
	%g5	(r05)		reserved for SPARC ABI
	%g6	(r06)		reserved for SPARC ABI
	%g7	(r07)		reserved for SPARC ABI
out	%o0	(r08)	[3]	outgoing parameter 0 / return value from callee
	%o1	(r09)	[1]	outgoing parameter 1
	%o2	(r10)	[1]	outgoing parameter 2
	%o3	(r11)	[1]	outgoing parameter 3
	%o4	(r12)	[1]	outgoing parameter 4
	%o5	(r13)	[1]	outgoing parameter 5
	%sp, %o6	(r14)	[1]	stack pointer
	%o7	(r15)	[1]	temporary value / address of CALL instruction
local	%l0	(r16)	[3]	local 0
	%l1	(r17)	[3]	local 1
	%l2	(r18)	[3]	local 2
	%l3	(r19)	[3]	local 3
	%l4	(r20)	[3]	local 4
	%l5	(r21)	[3]	local 5
	%l6	(r22)	[3]	local 6
	%l7	(r23)	[3]	local 7
in	%i0	(r24)	[3]	incoming parameter 0 / return value to caller
	%i1	(r25)	[3]	incoming parameter 1
	%i2	(r26)	[3]	incoming parameter 2
	%i3	(r27)	[3]	incoming parameter 3
	%i4	(r28)	[3]	incoming parameter 4
	%i5	(r29)	[3]	incoming parameter 5
	%fp, %i6	(r30)	[3]	frame pointer
	%i7	(r31)	[3]	return address - 8

Notes:

[1] assumed by caller to be destroyed (volatile) across a procedure call
 [2] should not be used by SPARC ABI library code
 [3] assumed by caller to be preserved across a procedure call

Figure 2 - Sparc register semantics

Particular compilers are likely to vary slightly.

Note that globals %g2-%g4 are reserved for the "application", which includes libraries and compiler. Thus, for example, libraries may overwrite these registers unless they've been compiled with suitable flags. Also, the "reserved" registers are presumed to be allocated (in the future) bottom-up, i.e. %g7 is currently the "safest" to use.

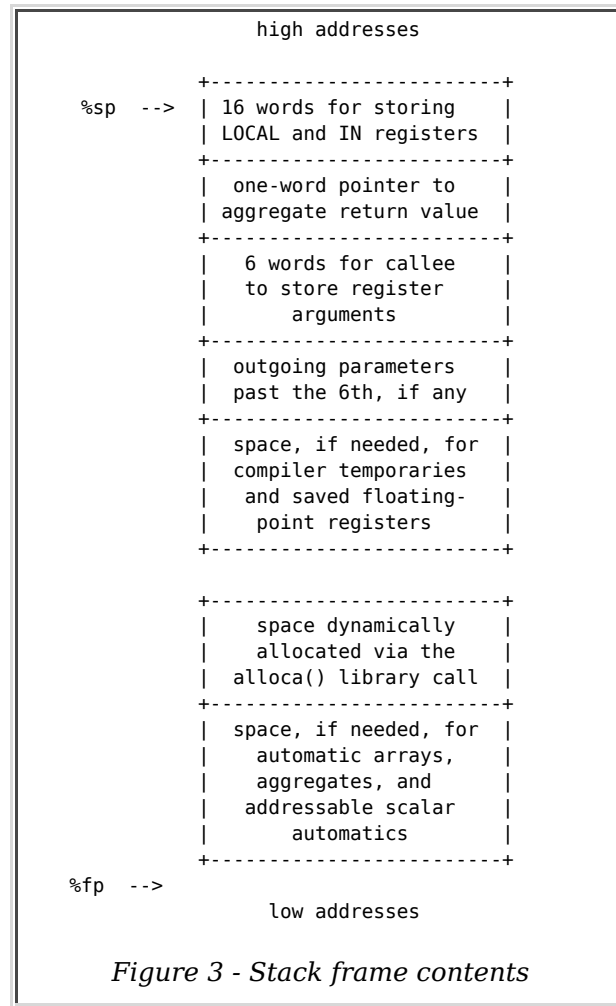
Optimizing linkers and interpreters are examples that use global registers.

Register Windows and the Stack

The sparc register windows are, naturally, intimately related to the stack. In particular, the stack pointer (%sp or %o6) must **always** point to a free block of 64 bytes. This area is used by the operating system (Solaris, SunOS, and Linux at least) to save the current *local* and *in* registers upon a system interrupt, exception, or trap instruction. (Note that this can occur at any time.)

Other aspects of register relations with memory are programming convention. The typical, and

recommended, layout of the stack is shown in figure 3. The figure shows a **stack frame**.



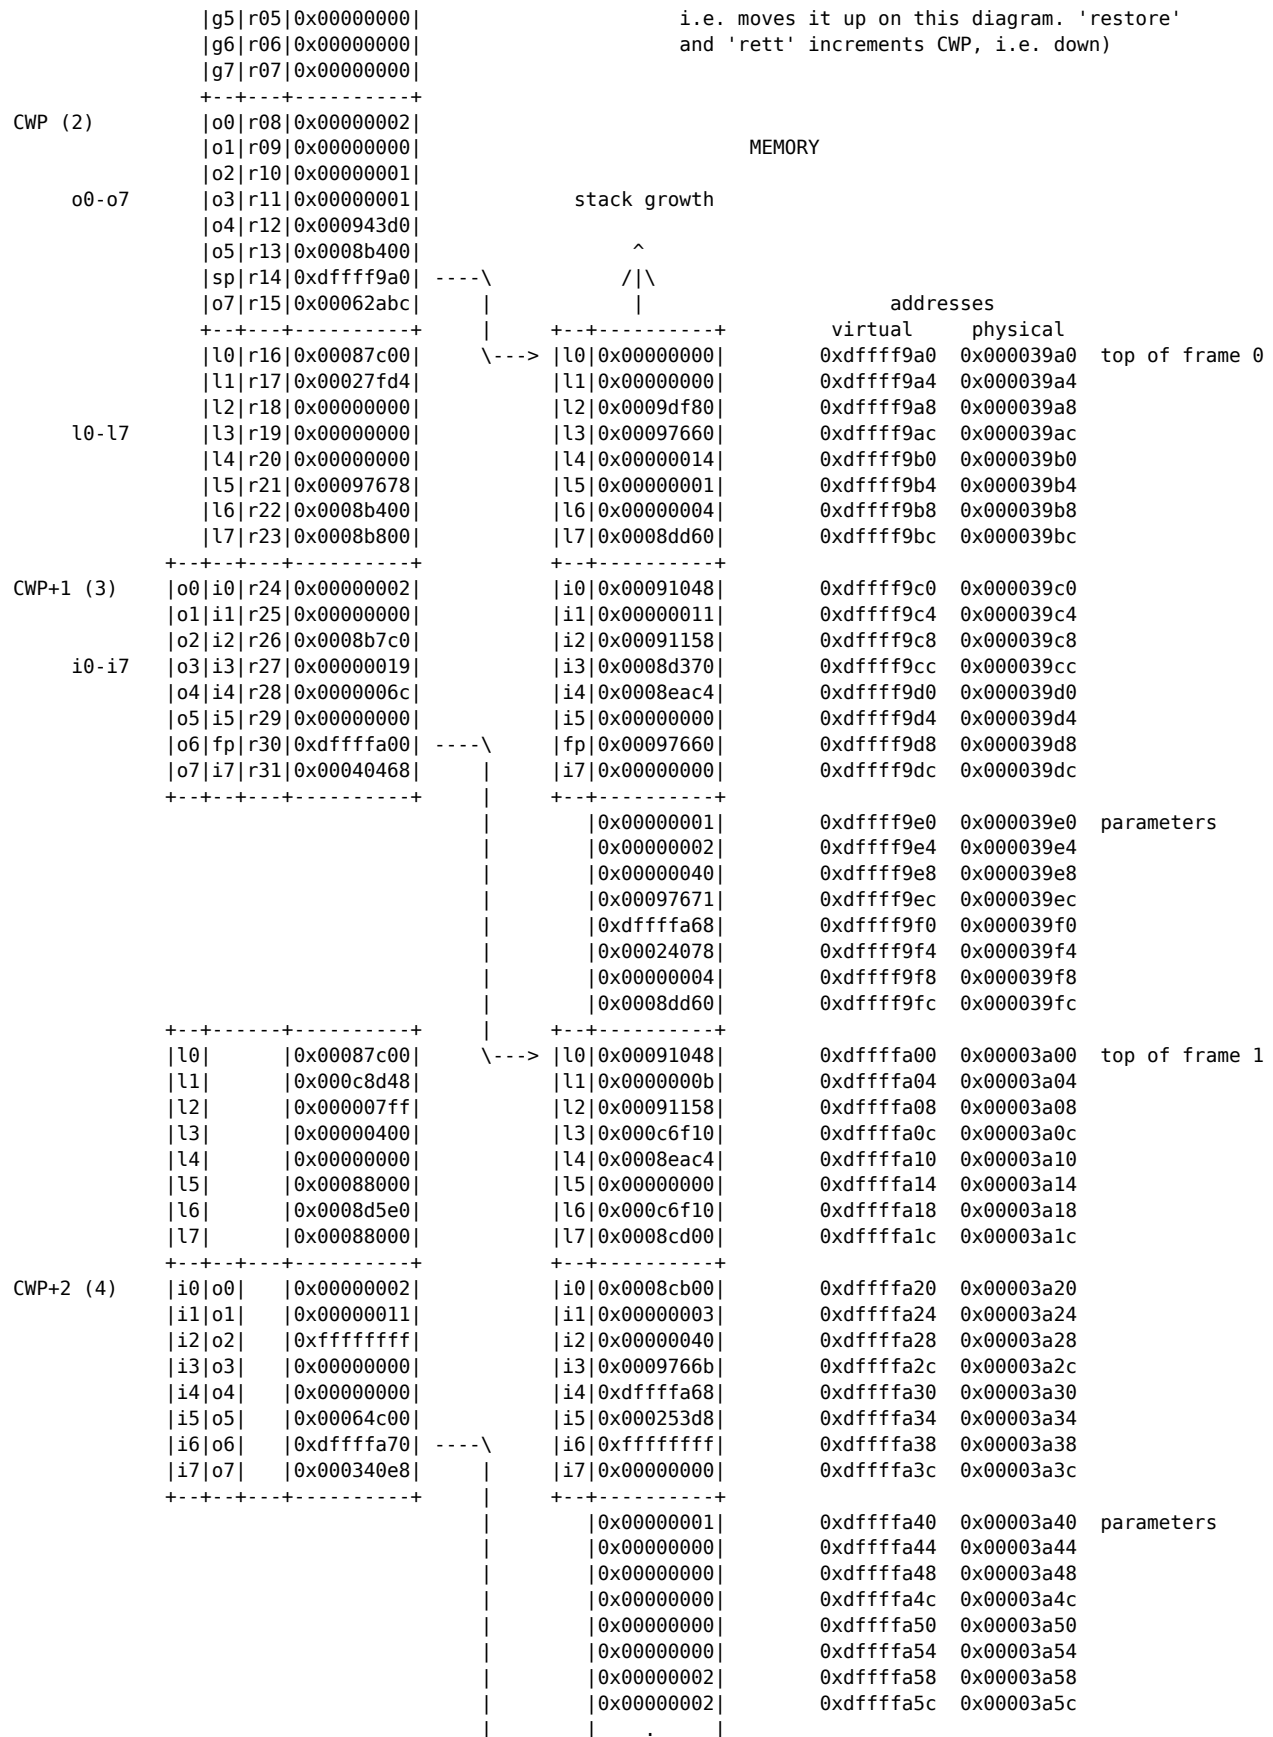
Note that the top boxes of figure 3 are addressed via the stack pointer (%sp), as positive offsets (including zero), and the bottom boxes are accessed over the frame pointer using negative offsets (excluding zero), and that the frame pointer is the old stack pointer. This scheme allows the separation of information known at compile time (number and size of local parameters, etc) from run-time information (size of blocks allocated by `alloca()`).

"addressable scalar automatics" is a fancy name for local variables.

The clever nature of the stack and frame pointers are that they are always 16 registers apart in the register windows. Thus, a SAVE instruction will make the current stack pointer into the frame pointer and, since the SAVE instruction also doubles as an ADD, create a new stack pointer. Figure 4 illustrates what the top of a stack might look like during execution. (The listing is from the "pwin" command in the [SimICS](#) simulator.)

REGISTER WINDOWS		
	+---+-----+	
	g0 r00 0x00000000	global
	g1 r01 0x00000006	registers
	g2 r02 0x00091278	
g0-g7	g3 r03 0x0008ebd0	
	g4 r04 0x00000000	

(note: 'save' and 'trap' decrements CWP,



				.. etc (another 16 bytes)

Figure 4 - Sample stack contents

Note how the stack contents are not necessarily synchronized with the registers. Various events can cause the register windows to be "flushed" to memory, including most system calls. A programmer can force this update by using `ST_FLUSH_WINDOWS` trap, which also reduces the number of valid windows to the minimum of 1.

Writing a library for multithreaded execution is an example that requires explicit flushing, as is `longjmp()`.

Procedure epilogue and prologue

The stack frame described in the previous section leads to the standard entry/exit mechanisms listed in figure 5.

```
function:
  save  %sp, -C, %sp

                                ; perform function, leave return value,
                                ; if any, in register %i0 upon exit

  ret   ; jmpl %i7+8, %g0
  restore ; restore %g0,%g0,%g0
```

Figure 5 - Epilogue/prologue in procedures

The `SAVE` instruction decrements the CWP, as discussed earlier, and also performs an addition. The constant "C" that is used in the figure to indicate the amount of space to make on the stack, and thus corresponds to the frame contents in Figure 3. The minimum is therefore the 16 words for the `LOCAL` and `IN` registers, i.e. (hex) 0x40 bytes.

A confusing element of the `SAVE` instruction is that the source operands (the first two parameters) are read from the old register window, and the destination operand (the rightmost parameter) is written to the new window. Thus, although "%sp" is indicated as both source and destination, the result is actually written into the stack pointer of the new window (the source stack pointer becomes renamed and is now the frame pointer).

The return instructions are also a bit particular. `ret` is a synthetic instruction, corresponding to `jmpl` (jump linked). This instruction jumps to the address resulting from adding 8 to the `%i7` register. The source instruction address (the address of the `ret` instruction itself) is written to the `%g0` register, i.e. it is discarded.

The `restore` instruction is similarly a synthetic instruction, and is just a short form for a `restore` that chooses not to perform an addition.

The calling instruction, in turn, typically looks as follows:

```
call <function>    ; jmpl <address>, %o7
mov 0, %o0
```

Again, the `call` instruction is synthetic, and is actually the same instruction that performs the return. This time, however, it is interested in saving the return address, into register `%o7`. Note that the delay slot is often filled with an instruction related to the parameters, in this example it

sets the first parameter to zero.

Note also that the return value is also generally passed in %o0.

Leaf procedures are different. A *leaf procedure* is an optimization that reduces unnecessary work by taking advantage of the knowledge that no `call` instructions exist in many procedures. Thus, the `save/restore` couple can be eliminated. The downside is that such a procedure may only use the *out* registers (since the *in* and *local* registers actually belong to the caller). See Figure 6.

```
function:
        ; no save instruction needed upon entry

        ; perform function, leave return value,
        ; if any, in register %o0 upon exit

    retl    ; jmpl %o7+8, %g0
    nop     ; the delay slot can be used for something else
```

Figure 6 - Epilogue/prologue in leaf procedures

Note in the figure that there is only one instruction overhead, namely the `retl` instruction. `retl` is also synthetic (return from leaf subroutine), is again a variant of the `jmpl` instruction, this time with %o7+8 as target.

Yet another variation of epilogue is caused by *tail call elimination*, an optimization supported by some compilers (including Sun's C compiler but not GCC). If the compiler detects that a called function will return to the calling function, it can replace its place on the stack with the called function. Figure 7 contains an example.

```
int
foo(int n)
{
    if (n == 0)
        return 0;
    else
        return bar(n);
}

    cmp     %o0,0
    bne     .L1
    or      %g0,%o7,%g1
    retl
    or      %g0,0,%o0
.L1:  call   bar
    or      %g0,%g1,%o7
```

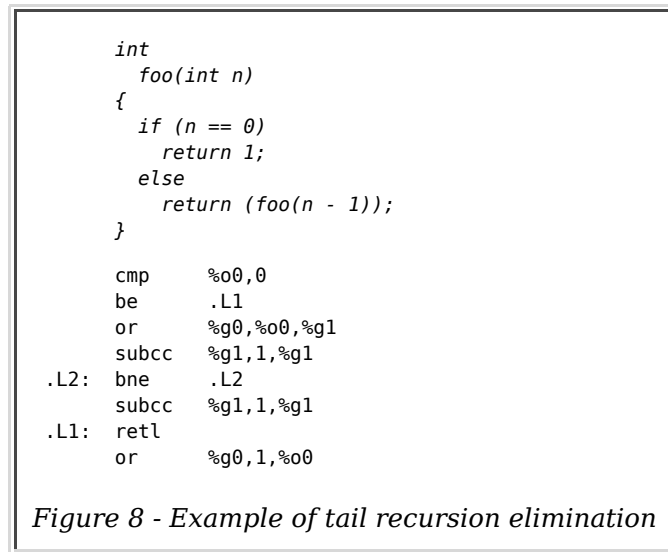
Figure 7 - Example of tail call elimination

Note that the `call` instruction overwrites register %o7 with the program counter. Therefore the above code saves the old value of %o7, and restores it in the delay slot of the `call` instruction. If the function call is register indirect, this twiddling with %o7 can be avoided, but of course that form of call is slower on modern processors.

The benefit of tail call elimination is to remove an indirection upon return. It is also needed to reduce register window usage, since otherwise the `foo()` function in Figure 7 would need to allocate a stack frame to save the program counter.

A special form of tail call elimination is *tail recursion elimination*, which detects functions calling

themselves, and replaces it with a simple branch. Figure 8 contains an example.



Needless to say, these optimizations produce code that is difficult to debug.

Procedures, stacks, and debuggers

When debugging an application, your debugger will be parsing the binary and consulting the symbol table to determine procedure entry points. It will also travel the stack frames "upward" to determine the current call chain.

When compiling for debugging, compilers will generate additional code as well as avoid some optimizations in order to allow reconstructing situations during execution. For example, GCC/GDB makes sure original parameter values are kept intact somewhere for future parsing of the procedure call stack. The live *in* registers other than %i0 are not touched. %i0 itself is copied into a free *local* register, and its location is noted in the symbol file. (You can find out where variables reside by using the "info address" command in GDB.)

Given that much of the semantics relating to stack handling and procedure call entry/exit code is only recommended, debuggers will sometimes be fooled. For example, the decision as to whether or not the current procedure is a leaf one or not can be incorrect. In this case a spurious procedure will be inserted between the current procedure and its "real" parent. Another example is when the application maintains its own implicit call hierarchy, such as jumping to function pointers. In this case the debugger can easily become totally confused.

The window overflow and underflow traps

When the SAVE instruction decrements the current window pointer (CWP) so that it coincides with the invalid window in the window invalid mask (WIM), a **window overflow** trap occurs. Conversely, when the RESTORE or RETT instructions increment the CWP to coincide with the invalid window, a **window underflow** trap occurs.

Either trap is handled by the operating system. Generally, data is written out to memory and/or read from memory, and the WIM register suitably altered.

The code in Figure 9 and Figure 10 below are bare-bones handlers for the two traps. The text is directly from the source code, and sort of works. (As far as I know, these are minimalistic handlers

for Sparc V8). Note that there is no way to directly access window registers other than the current one, hence the code does additional save/restore instructions. It's pretty tricky to understand the code, but figure 1 should be of help.

```

/* a SAVE instruction caused a trap */
window_overflow:
/* rotate WIM on bit right, we have 8 windows */
mov %wim,%l3
sll %l3,7,%l4
srl %l3,1,%l3
or %l3,%l4,%l3
and %l3,0xff,%l3

/* disable WIM traps */
mov %g0,%wim
nop; nop; nop

/* point to correct window */
save

/* dump registers to stack */
std %l0, [%sp + 0]
std %l2, [%sp + 8]
std %l4, [%sp + 16]
std %l6, [%sp + 24]
std %i0, [%sp + 32]
std %i2, [%sp + 40]
std %i4, [%sp + 48]
std %i6, [%sp + 56]

/* back to where we should be */
restore

/* set new value of window */
mov %l3,%wim
nop; nop; nop

/* go home */
jmp %l1
rett %l2

```

Figure 9 - window_underflow trap handler

```

/* a RESTORE instruction caused a trap */
window_underflow:

/* rotate WIM on bit LEFT, we have 8 windows */
mov %wim,%l3
srl %l3,7,%l4
sll %l3,1,%l3
or %l3,%l4,%l3
and %l3,0xff,%l3

/* disable WIM traps */
mov %g0,%wim
nop; nop; nop

/* point to correct window */
restore
restore

/* dump registers to stack */
ldd [%sp + 0], %l0
ldd [%sp + 8], %l2

```

```
    ldd [%sp + 16], %l4
    ldd [%sp + 24], %l6
    ldd [%sp + 32], %i0
    ldd [%sp + 40], %i2
    ldd [%sp + 48], %i4
    ldd [%sp + 56], %i6

    /* back to where we should be */
    save
    save

    /* set new value of window */
    mov %l3,%wim
    nop; nop; nop

    /* go home */
    jmp %l1
    rett %l2
```

Figure 10 - window_underflow trap handler

Note: some of the figures and data is (c) copyright Sun Microsystems. I can't imagine they would object to my usage of the material, but if you make copies you are hereby advised.

Created and maintained by [Peter Magnusson](#).

Created in March 1997, last revision in April 1997.